

# Sparse Warnings

Daniel Axtens

`linux.conf.au` 2017 Kernel Miniconf

```
make C=2 ...
```

# Where we're going

- 1 Introduction
- 2 What does Sparse detect?
  - Simple things
  - Enhancing sparse with type annotations
- 3 How is Sparse actually used?
- 4 Can we improve the situation?
  - Long term
  - Short term
- 5 Where to from here
  - Sparse as a gateway to kernel development
- 6 Conclusion

# Simple static analysis

- Undefined/unwise behaviour:

warning: expression using sizeof bool

warning: odd constant \_Bool cast (ffffffffffffffff becomes 1)

- Odd accesses:

warning: invalid access past the end of 's32' (12 8)

- Static suggestions:

warning: symbol 'ppc\_fadvise64\_64' was not declared. \  
Should it be static?

# What does sparse understand?

sparse + annotations  $\Rightarrow$  understanding more than the C compiler alone:

- Endianness of variables
- Address space of pointers
- Pointers that should not be dereferenced
- Types needing explicit conversion
- and (probably) more...

# Base types

```
unsigned int instr;  
instr = cpu_to_le32(instr);
```

# Base types

```
unsigned int instr;  
instr = cpu_to_le32(instr);
```

---

```
warning: incorrect type in assignment (different base types)  
    expected unsigned int [unsigned] [assigned] instr  
    got restricted __le32 [usertype] <noident>
```

# Address spaces

```
unsigned long pc;  
int instr;
```

```
probe_kernel_address((unsigned int __user *)pc, instr);
```



# Address spaces

```
unsigned long pc;  
int instr;
```

```
probe_kernel_address((unsigned int __user *)pc, instr);
```

---

```
warning: incorrect type in argument 2 (different address spaces)  
    expected void const *src  
    got unsigned int [noderef] <asn:1>*<noident>
```

# Address spaces

```
struct rt_sigframe __user *frame;  
  
printk_ratelimited(regs->msr & MSR_64BIT ? fmt64 : fmt32,  
                  tsk->comm, tsk->pid, "setup_rt_frame",  
                  (long)frame, regs->nip, regs->link);
```

# Address spaces

```
struct rt_sigframe __user *frame;

printk_ratelimited(regs->msr & MSR_64BIT ? fmt64 : fmt32,
                  tsk->comm, tsk->pid, "setup_rt_frame",
                  (long)frame, regs->nip, regs->link);
```

---

warning: cast removes address space of expression

# But wait, there's more!

- no cast types

```
arch/powerpc/kernel/time.c:361:37: warning: implicit cast to nocrash type  
arch/powerpc/kernel/time.c:362:29: warning: implicit cast to nocrash type
```

- no dereference pointers (e.g. IO)

```
arch/powerpc/kernel/io.c:40:24: warning: dereference of noderef expression  
arch/powerpc/kernel/io.c:56:18: warning: dereference of noderef expression
```

- restricted types

```
arch/powerpc/sysdev/mpic.c:356:18: warning: cast to restricted __le32
```

- 1 Introduction
- 2 What does Sparse detect?
  - Simple things
  - Enhancing sparse with type annotations
- 3 How is Sparse actually used?
- 4 Can we improve the situation?
  - Long term
  - Short term
- 5 Where to from here
  - Sparse as a gateway to kernel development
- 6 Conclusion




Figure: Sparse warnings, 2 ppc defconfigs, 2016

- 1 Introduction
- 2 What does Sparse detect?
  - Simple things
  - Enhancing sparse with type annotations
- 3 How is Sparse actually used?
- 4 Can we improve the situation?
  - Long term
  - Short term
- 5 Where to from here
  - Sparse as a gateway to kernel development
- 6 Conclusion

# In the long term...

Fix the warnings



# PowerPC warnings




Figure: arch/powerpc sparse warnings, combination of 2 defconfigs, 2016

# PowerPC warnings





Figure: arch/powerpc sparse warnings, combination of 2 defconfigs, 2016-present

# Total warnings



## linux/types.h: enable endian checks for all sparse builds

By now, linux is mostly endian-clean. Enabling endian-ness checks for everyone produces about 200 new sparse warnings for me - less than 10% over the 2000 sparse warnings already there.

Not a big deal, OTOH enabling this helps people notice they are introducing new bugs.

So let's just drop `__CHECK_ENDIAN__`. Follow-up patches can drop distinction between `__bitwise` and `__bitwise__`.

Cc: Linus Torvalds <torvalds@linux-foundation.org>

Suggested-by: Christoph Hellwig <hch@infradead.org>

Signed-off-by: Michael S. Tsirkin <mst@redhat.com>

## diff <before> <after>

arch/powerpc/kernel/nvram\_64.c:1177:32: warning: cast to restricted \_\_be16

arch/powerpc/kernel/nvram\_64.c:893:22: warning: incorrect type in assignment \  
(different base types)  
expected unsigned short [unsigned] [addressable] length  
got restricted \_\_be16 [usertype] <noident>

arch/powerpc/kvm/book3s\_64\_vio\_hv.c:282:37: warning: cast to restricted \_\_be64

arch/powerpc/kvm/book3s\_hv\_builtin.c:421:22: warning: incorrect type in assignment \  
(different base types)  
expected restricted \_\_be32 [addressable] [usertype] xirr  
got unsigned int

arch/powerpc/perf/hv-24x7.c:1166:18: warning: cast to restricted \_\_be64

# In the long term...

Fix the warnings

## How do you diff compiler warnings/sparse output?

- `grep 'arch/powerpc' sparse-output | wc -l`
  - Simple, low fidelity.
- `diff sparse-output-1 sparse-output-2`
  - Reordering due to parallel builds
  - Changing line numbers
- Write our own.

# Introducing smart-sparse-diff

<https://github.com/daxtens/smart-sparse-diff>



# Where to from here

- **Fix sparse warnings in your code**
- When they have reached an acceptably low quantity
  - Contributors: don't add warnings
  - Maintainers: require it of contributors
  - ML reviewers: evaluate it in your reviews

# Sparse as a gateway to kernel development

- Pick a file with sparse warnings
- How do patches go into that file?
  - What mailing list?
  - What subject line format?
- When doing patches:
  - Don't fix just one of many, fix all of one type
  - Compile (and if appropriate, run) test!
- Consult docs on submitting patches:
  - Commit messages
  - How to send a non-broken email
- Expect bike-shedding.

# Conclusion

- 1 Introduction
- 2 What does Sparse detect?
  - Simple things
  - Enhancing sparse with type annotations
- 3 How is Sparse actually used?
- 4 Can we improve the situation?
  - Long term
  - Short term
- 5 Where to from here
  - Sparse as a gateway to kernel development
- 6 Conclusion

- This work represents the view of the author and does not necessarily represent the view of IBM.
- IBM is a registered trademark of International Business Machines Corporation in the United States and/or other countries.
- Linux is a registered trademark of Linus Torvalds.
- Other company, product, and service names may be trademarks or service marks of others.